
Support Vector Machines for Credit Default Prediction

Matthew Collins
Department of Computer Science
Clemson University
mbc3@clemson.edu

Abstract

1 This class project presents an analysis of support vector machines (SVMs) applied
2 to the UCI Default of Credit Card Clients dataset. The project explores the impacts
3 of different kernel functions and optimization techniques on classification perfor-
4 mance. SVMs with linear, radial basis function (RBF), polynomial, and sigmoid
5 kernels are evaluated. A grid search approach is used for hyperparameter optimiza-
6 tion, and the decision boundary is visualized using Principal Component Analysis
7 (PCA). The results are evaluated using accuracy, precision, recall, F1-score, and
8 runtime, demonstrating the utility of SVMs for binary classification tasks on tabular
9 data.

10 1 Introduction

11 Support vector machines are a powerful supervised machine learning approach that have demonstrated
12 robust performance in a wide range of classification tasks. This project investigates the application of
13 SVMs to the UCI Default of Credit Card Clients dataset, with a specific focus on assessing the impact
14 of kernel functions and optimization techniques on model performance. The study also examines
15 how hyperparameter tuning can enhance SVM efficiency and accuracy, while PCA is employed to
16 visualize the decision boundary for improved interpretability. The primary objective of this study is
17 to assess the efficacy of SVMs for predicting credit card default and to analyze the computational
18 trade offs associated with various kernel functions and optimization strategies.

19 2 Dataset

20 The UCI Default of Credit Card Clients dataset is a multivariate dataset containing data from 30,000
21 credit card holders in Taiwan. The dataset is primarily designed for binary classification tasks,
22 specifically predicting whether a client would default on their next month's payment. It includes
23 23 features, which consist of demographic variables such as age, sex, education, and marital status,
24 along with financial metrics like credit limits, historical bill statements, and repayment amounts. The
25 target variable is binary, indicating whether a client defaulted (1) or did not default (0) on their credit
26 card payment.

27 The dataset includes a mix of categorical and numerical variables. Key variables include repayment
28 statuses (PAY_0 through PAY_6), historical bill amounts (BILL_AMT1 through BILL_AMT6), and
29 historical repayment amounts (PAY_AMT1 through PAY_AMT6). Additionally, demographic variables
30 such as SEX, EDUCATION, and MARRIAGE provide insights into customer profiles. Table 1 summarizes
31 the features included in the dataset.

Table 1: Summary of Key Features in the Dataset

Variable Name	Type	Description	Units
SEX	Categorical	Gender	1 = Male, 2 = Female
EDUCATION	Categorical	Education level	1 = Graduate, etc.
MARRIAGE	Categorical	Marital status	1 = Married, etc.
BILL_AMT1-6	Numerical	Monthly bill amounts (6 months)	NT Dollars
PAY_AMT1-6	Numerical	Historical repayment amounts (6 months)	NT Dollars
PAY_0-6	Categorical	Repayment status (-2: advance payment)	Discrete values
DEFAULT (Target)	Binary	Default status for the next month	0 = No, 1 = Yes

To prepare the dataset for modeling, extensive preprocessing was conducted. Several variables, such as SEX, EDUCATION, MARRIAGE, and repayment statuses (PAY_0 through PAY_6), were categorical in nature. Since support vector machines require numerical input, these categorical variables were transformed into one hot encoded binary variables. This encoding ensured that each category was treated as an independent feature, avoiding implicit ordinal assumptions. Furthermore, the dataset exhibited significant class imbalance, with the majority of clients not defaulting on their payments. To address this imbalance, the majority class (non-default clients) was downsampled to create a balanced dataset, ensuring that the model learned patterns equally well from both defaulting and non-defaulting clients.

Another essential preprocessing step involved standardizing numerical features, such as bill amounts and repayment amounts. These features had varying scales, with some variables measured in monetary units and others as categorical values. Without standardization, features with larger magnitudes could disproportionately influence the model. Therefore, all continuous features were scaled to have a mean of zero and a standard deviation of one, ensuring fair contribution from each variable.

The preprocessing steps of one hot encoding, downsampling, and standardization were necessary to ensure that the dataset was compatible with support vector machines and capable of producing meaningful predictions. These transformations addressed the challenges posed by mixed feature types, class imbalance, and variable scaling, ultimately improving the efficiency and accuracy of the machine learning models applied in this study.

3 Methodology

After preprocessing the dataset to address class imbalance and ensure numerical compatibility, the project proceeded with training a preliminary Support Vector Machine (SVM) using the radial basis function (RBF) kernel. The choice of the RBF kernel for initial experimentation was motivated by its capability to map input features into higher-dimensional spaces, enabling it to capture complex, nonlinear decision boundaries. The initial model utilized default hyperparameters, with the regularization parameter $C = 1$ and kernel coefficient $\gamma = \frac{1}{\text{number of features}}$, to establish a baseline performance. Model evaluation involved generating confusion matrices, as shown in Figure 1, and calculating accuracy scores.

The SVM optimization process centered around hyperparameter tuning using grid search cross-validation. The grid search method systematically evaluates combinations of C and γ values to identify the optimal hyperparameters that maximize classification performance. The regularization parameter C controls the trade off between achieving a low training error and maintaining a large margin. The SVM optimization problem is expressed as:

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i,$$

subject to:

$$y_i(\mathbf{w}^\top \phi(\mathbf{x}_i) + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, n,$$

where \mathbf{w} is the weight vector, b is the bias, ξ_i are slack variables allowing soft margin violations, and $\phi(\mathbf{x}_i)$ represents the transformation of input features into a higher-dimensional space. The kernel coefficient γ determines the influence of individual training points, effectively controlling the curvature of the decision boundary in the feature space. The post-optimization confusion matrix, which demonstrates improved performance, is shown in Figure 2.

To provide a comprehensive analysis, four kernel functions—linear, polynomial, RBF, and sigmoid—were evaluated. Each kernel was optimized using grid search to determine its best performing hyperparameters. The linear kernel assumes a straight-line decision boundary and is suitable for datasets that are linearly separable.

Performance metrics, including accuracy, precision, recall, F1-score, and runtime, were recorded for each kernel to evaluate their suitability for the credit card default prediction task. The accuracy of the SVM models was calculated as:

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}},$$

where TP, TN, FP, and FN represent true positives, true negatives, false positives, and false negatives, respectively. Precision, recall, and F1-score were also computed to provide a more detailed assessment of the models' performance in handling class imbalance.

To facilitate visualization and interpretability, Principal Component Analysis (PCA) was employed to reduce the feature space to two dimensions. PCA identifies orthogonal directions, or principal components, that capture the maximum variance in the dataset. The decision boundary for each kernel, projected onto the first two principal components, is illustrated in Figure 4. The PCA transformed decision boundary provides a visual representation of the separability between defaulting and non-defaulting clients.

4 Results

The comparative analysis of kernel functions yielded results, summarized in Table 2. The RBF kernel outperformed other kernels, achieving the highest accuracy (68.75%) and F1-score (0.65). The scree plot in Figure 3 highlights the explained variance by the principal components, providing insights into the dimensionality reduction process.




Figure 1: Confusion Matrix Before Grid Search.




Figure 2: Confusion Matrix After Grid Search.

Table 2: Kernel Comparison Results

Kernel	Accuracy	Precision	Recall	F1-Score	Runtime (s)	Best Hyperparameters (C, γ)
Linear	0.6525	0.7095	0.5224	0.6017	0.2457	(1, Scale)
Polynomial	0.6700	0.7226	0.5572	0.6292	0.0629	(10, 0.01)
RBF	0.6875	0.7436	0.5771	0.6499	0.0634	(1, 0.01)
Sigmoid	0.6450	0.7007	0.5124	0.5920	0.0891	(100, 0.001)




Figure 3: Scree Plot Showing Percentage of Explained Variance.




Figure 4: Decision Surface Using PCA Transformed Features.

92 5 Discussion

93 The results highlight the importance of kernel selection in Support Vector Machines (SVMs) for binary
 94 classification tasks. Among the evaluated kernels, the radial basis function (RBF) kernel emerged as
 95 the best performing option, achieving an accuracy of 68.75%. This performance underscores the RBF
 96 kernel’s strength in mapping data into higher-dimensional spaces, enabling it to capture nonlinear
 97 decision boundaries. However, the achieved accuracy is relatively modest, raising questions about
 98 the suitability of SVMs for this particular dataset. While the polynomial kernel showed competitive
 99 performance with an accuracy of 67.00%, its sensitivity to parameter tuning limited its robustness.
 100 The linear kernel underperformed, achieving an accuracy of only 65.25%, which can be attributed
 101 to its inability to handle the nonlinear structure of the data. The sigmoid kernel, with its complex
 102 parameterization, struggled to converge effectively and yielded the lowest accuracy of 64.50%.

103 The relatively low accuracy across all kernels suggests that SVMs, while powerful, may not be
 104 the optimal model for this task. The imbalanced and high-dimensional nature of the dataset likely
 105 contributed to the suboptimal performance. Although hyperparameter tuning via grid search improved
 106 the models, as evidenced by the enhanced confusion matrix after optimization (Figure 2), the
 107 gains were incremental and insufficient to achieve high accuracy. This indicates that the inherent
 108 characteristics of the dataset—such as overlapping class distributions and complex relationships
 109 among features—pose challenges that SVMs alone may not effectively address.

110 The PCA visualization of the decision boundary (Figure 4) provided insights into the separability
 111 of the classes, but the limited explained variance from the first two principal components (Figure 3)
 112 suggests that the majority of the dataset’s complexity lies in higher dimensions. This highlights
 113 a limitation of using PCA for interpretability and suggests that other dimensionality reduction
 114 techniques, could provide more nuanced insights into the data’s structure.

6 Conclusion

This study demonstrates the effectiveness of support vector machines for credit card default prediction. The findings emphasize the importance of kernel selection, with the RBF kernel offering the best performance among the four kernels evaluated. Hyperparameter tuning further enhanced classification accuracy, though its impact was secondary to the choice of kernel. Future work could explore additional dimensionality reduction techniques, alternative kernel functions, and the application of SVMs to larger, more diverse datasets. In summary, while SVMs demonstrated some utility in this study, the results indicate that they may not be the most effective model for predicting credit card default. The modest accuracy, despite careful preprocessing and hyperparameter tuning, underscores the need for exploring alternative machine learning models and techniques to better address the complexity of the dataset and improve predictive performance.

References

- Chang, C. and Lin, C. (2011). LIBSVM: A library for support vector machines. *ACM Transactions on Intelligent Systems and Technology*, 2(3), 27.
- UCI Machine Learning Repository: Default of Credit Card Clients Dataset. Available at <https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>.
- Cortes, C., and Vapnik, V. (1995). Support-Vector Networks. *Machine Learning*, 20(3), 273–297.