
Support Vector Machines for Credit Default Prediction

Matthew Collins
Department of Computer Science
Clemson University
mbc3@clemson.edu

Abstract

1 This class project presents an analysis of support vector machines (SVMs) applied
2 to the UCI Default of Credit Card Clients dataset. The project explores the impacts
3 of different kernel functions and optimization techniques on classification perfor-
4 mance. SVMs with linear, radial basis function (RBF), polynomial, and sigmoid
5 kernels are evaluated. A grid search approach is used for hyperparameter optimiza-
6 tion, and the decision boundary is visualized using Principal Component Analysis
7 (PCA). The results are evaluated using accuracy, precision, recall, F1-score, and
8 runtime, demonstrating the utility of SVMs for binary classification tasks on tabular
9 data.

10 1 Introduction

11 Support vector machines are a powerful supervised machine learning approach that have demonstrated
12 robust performance in a wide range of classification tasks. This project investigates the application of
13 SVMs to the UCI Default of Credit Card Clients dataset, with a specific focus on assessing the impact
14 of kernel functions and optimization techniques on model performance. The study also examines
15 how hyperparameter tuning can enhance SVM efficiency and accuracy, while PCA is employed to
16 visualize the decision boundary for improved interpretability. The primary objective of this study is
17 to assess the efficacy of SVMs for predicting credit card default and to analyze the computational
18 trade offs associated with various kernel functions and optimization strategies.

19 2 Dataset

20 The UCI Default of Credit Card Clients dataset is a multivariate dataset containing data from 30,000
21 credit card holders in Taiwan. The dataset is primarily designed for binary classification tasks,
22 specifically predicting whether a client would default on their next month's payment. It includes
23 23 features, which consist of demographic variables such as age, sex, education, and marital status,
24 along with financial metrics like credit limits, historical bill statements, and repayment amounts. The
25 target variable is binary, indicating whether a client defaulted (1) or did not default (0) on their credit
26 card payment.

27 The dataset includes a mix of categorical and numerical variables. Key variables include repayment
28 statuses (PAY_0 through PAY_6), historical bill amounts (BILL_AMT1 through BILL_AMT6), and
29 historical repayment amounts (PAY_AMT1 through PAY_AMT6). Additionally, demographic variables
30 such as SEX, EDUCATION, and MARRIAGE provide insights into customer profiles. Table 1 summarizes
31 the features included in the dataset.

Table 1: Summary of Key Features in the Dataset

Variable Name	Type	Description	Units
SEX	Categorical	Gender	1 = Male, 2 = Female
EDUCATION	Categorical	Education level	1 = Graduate, etc.
MARRIAGE	Categorical	Marital status	1 = Married, etc.
BILL_AMT1-6	Numerical	Monthly bill amounts (6 months)	NT Dollars
PAY_AMT1-6	Numerical	Historical repayment amounts (6 months)	NT Dollars
PAY_0-6	Categorical	Repayment status (-2: advance payment)	Discrete values
DEFAULT (Target)	Binary	Default status for the next month	0 = No, 1 = Yes

To prepare the dataset for modeling, extensive preprocessing was conducted. Several variables, such as SEX, EDUCATION, MARRIAGE, and repayment statuses (PAY_0 through PAY_6), were categorical in nature. Since support vector machines require numerical input, these categorical variables were transformed into one hot encoded binary variables. This encoding ensured that each category was treated as an independent feature, avoiding implicit ordinal assumptions. Furthermore, the dataset exhibited significant class imbalance, with the majority of clients not defaulting on their payments. To address this imbalance, the majority class (non-default clients) was downsampled to create a balanced dataset, ensuring that the model learned patterns equally well from both defaulting and non-defaulting clients.

Another essential preprocessing step involved standardizing numerical features, such as bill amounts and repayment amounts. These features had varying scales, with some variables measured in monetary units and others as categorical values. Without standardization, features with larger magnitudes could disproportionately influence the model. Therefore, all continuous features were scaled to have a mean of zero and a standard deviation of one, ensuring fair contribution from each variable.

The preprocessing steps of one hot encoding, downsampling, and standardization were necessary to ensure that the dataset was compatible with support vector machines and capable of producing meaningful predictions. These transformations addressed the challenges posed by mixed feature types, class imbalance, and variable scaling, ultimately improving the efficiency and accuracy of the machine learning models applied in this study.

3 Methodology

The Support Vector Machine (SVM) algorithm aims to find an optimal hyperplane that maximizes the margin between two classes while minimizing classification errors. For a training dataset with n points $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ where $\mathbf{x}_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$, the primal optimization problem is formulated as:

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

subject to the constraints:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, \dots, n$$

where \mathbf{w} is the normal vector to the hyperplane, b is the bias term, ξ_i are slack variables allowing for soft margin violations, and C is the regularization parameter controlling the trade-off between margin maximization and error minimization.

The dual formulation of this optimization problem, derived using Lagrange multipliers, is:

$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j$$

61 subject to:

$$0 \leq \alpha_i \leq C, \quad i = 1, \dots, n \quad \text{and} \quad \sum_{i=1}^n \alpha_i y_i = 0$$

62 To handle nonlinear classification problems, kernel functions $K(\mathbf{x}_i, \mathbf{x}_j)$ are introduced to implicitly map the input features into a higher-dimensional space. The kernel trick transforms the dual optimization problem to:

$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$$

65 In this project, four kernel functions were evaluated:

- 66 1. Linear: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \mathbf{x}_j$
- 67 2. Polynomial: $K(\mathbf{x}_i, \mathbf{x}_j) = (\gamma \mathbf{x}_i^\top \mathbf{x}_j + r)^d$
- 68 3. RBF (Gaussian): $K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2)$
- 69 4. Sigmoid: $K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\gamma \mathbf{x}_i^\top \mathbf{x}_j + r)$

70 The decision function for classifying new points becomes:

$$f(\mathbf{x}) = \text{sign} \left(\sum_{i=1}^n \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b \right)$$

71 Grid search cross-validation was employed to optimize the hyperparameters C and γ for each kernel.
72 The search space included:

- 73 • $C \in \{0.1, 1, 10, 100\}$
- 74 • $\gamma \in \{0.001, 0.01, 0.1, 1\}$

75 To facilitate visualization and interpretation, Principal Component Analysis (PCA) was applied to
76 reduce the feature space to two dimensions. Given the centered data matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$, where n is
77 the number of samples and d is the number of features, PCA first computes the covariance matrix:

$$\Sigma = \frac{1}{n} \mathbf{X}^\top \mathbf{X}$$

78 PCA then performs eigendecomposition of the covariance matrix:

$$\Sigma = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^\top$$

79 where $\mathbf{V} \in \mathbb{R}^{d \times d}$ contains the eigenvectors as columns (sorted by decreasing eigenvalues) and
80 $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_d)$ contains the corresponding eigenvalues. The transformed features are obtained
81 by projecting the data onto the first two principal components:

$$\mathbf{X}_{\text{transformed}} = \mathbf{X} \mathbf{V}_{1:2}$$

82 where $\mathbf{V}_{1:2} \in \mathbb{R}^{d \times 2}$ contains the first two eigenvectors. This projection preserves the maximum
83 possible variance in the data while reducing dimensionality to facilitate visualization of the decision
84 boundaries.

4 Results

The comparative analysis of kernel functions yielded results, summarized in Table 2. The RBF kernel outperformed other kernels, achieving the highest accuracy (68.75%) and F1-score (0.65). The scree plot in Figure 3 highlights the explained variance by the principal components, providing insights into the dimensionality reduction process.




Figure 1: Confusion Matrix Before Grid Search.

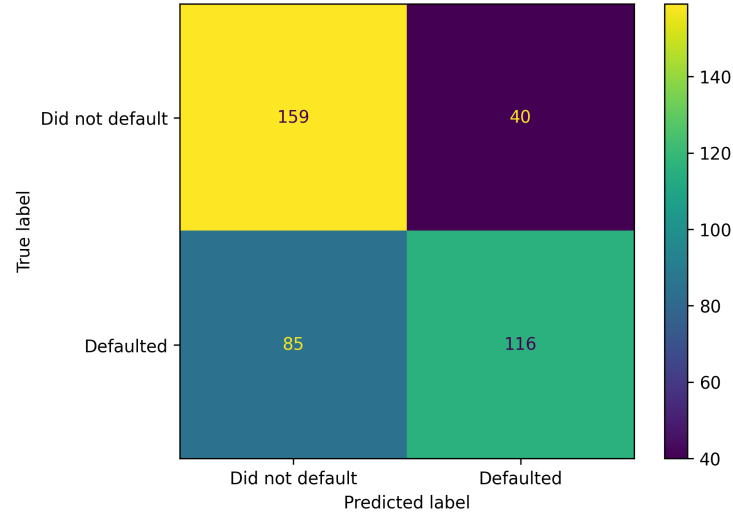


Figure 2: Confusion Matrix After Grid Search.

Table 2: Kernel Comparison Results

Kernel	Accuracy	Precision	Recall	F1-Score	Runtime (s)	Best Hyperparameters (C, γ)
Linear	0.6525	0.7095	0.5224	0.6017	0.2457	(1, Scale)
Polynomial	0.6700	0.7226	0.5572	0.6292	0.0629	(10, 0.01)
RBF	0.6875	0.7436	0.5771	0.6499	0.0634	(1, 0.01)
Sigmoid	0.6450	0.7007	0.5124	0.5920	0.0891	(100, 0.001)




Figure 3: Scree Plot Showing Percentage of Explained Variance.




Figure 4: Decision Surface Using PCA Transformed Features.

90 5 Discussion

91 The results highlight the importance of kernel selection in Support Vector Machines (SVMs) for binary
 92 classification tasks. Among the evaluated kernels, the radial basis function (RBF) kernel emerged as
 93 the best performing option, achieving an accuracy of 68.75%. This performance underscores the RBF
 94 kernel's strength in mapping data into higher-dimensional spaces, enabling it to capture nonlinear

95 decision boundaries. However, the achieved accuracy is relatively modest, raising questions about
96 the suitability of SVMs for this particular dataset. While the polynomial kernel showed competitive
97 performance with an accuracy of 67.00%, its sensitivity to parameter tuning limited its robustness.
98 The linear kernel underperformed, achieving an accuracy of only 65.25%, which can be attributed
99 to its inability to handle the nonlinear structure of the data. The sigmoid kernel, with its complex
100 parameterization, struggled to converge effectively and yielded the lowest accuracy of 64.50%.

101 The relatively low accuracy across all kernels suggests that SVMs, while powerful, may not be
102 the optimal model for this task. The imbalanced and high-dimensional nature of the dataset likely
103 contributed to the suboptimal performance. Although hyperparameter tuning via grid search improved
104 the models, as evidenced by the enhanced confusion matrix after optimization (Figure 2), the
105 gains were incremental and insufficient to achieve high accuracy. This indicates that the inherent
106 characteristics of the dataset—such as overlapping class distributions and complex relationships
107 among features—pose challenges that SVMs alone may not effectively address.

108 The PCA visualization of the decision boundary (Figure 4) provided insights into the separability
109 of the classes, but the limited explained variance from the first two principal components (Figure 3)
110 suggests that the majority of the dataset's complexity lies in higher dimensions. This highlights
111 a limitation of using PCA for interpretability and suggests that other dimensionality reduction
112 techniques, could provide more nuanced insights into the data's structure.

113 6 Conclusion

114 This project reveals the limitations of support vector machines for credit card default prediction, with
115 even the best performing RBF kernel achieving only 68.75% accuracy. While the findings emphasize
116 the importance of kernel selection, none of the evaluated kernels achieved satisfactory performance
117 for real-world application. Although hyperparameter tuning through grid search provided some
118 improvements, the gains were minimal and failed to address the fundamental limitations of SVMs for
119 this particular problem.

120 The consistently suboptimal performance across all kernel types, ranging from 64.50% to 68.75%
121 accuracy, suggests that SVMs may not be well-suited for credit card default prediction on this
122 dataset. The complex, high-dimensional nature of the data, combined with likely class overlap, poses
123 significant challenges that SVMs struggle to overcome. Future work should focus on exploring more
124 sophisticated machine learning approaches, such as ensemble methods or deep learning models, which
125 may be better equipped to handle the intricacies of credit default prediction. Additionally, investigating
126 alternative feature engineering techniques and addressing the dataset's inherent complexities could
127 prove more beneficial than further optimization of SVM models.

128 References

- 129 • Agarwal, A. (2023). Everything About SVM Classification: Above and Be-
130 yond. *Towards Data Science*. Available at [https://towardsdatascience.com/](https://towardsdatascience.com/everything-about-svm-classification-above-and-beyond-cc665bfd993e)
131 [everything-about-svm-classification-above-and-beyond-cc665bfd993e](https://towardsdatascience.com/everything-about-svm-classification-above-and-beyond-cc665bfd993e).
- 132 • Bhatia, N. (2023). Support Vector Machines (SVM): An Intuitive Explanation. *Medium*.
133 Available at [https://medium.com/low-code-for-advanced-data-science/](https://medium.com/low-code-for-advanced-data-science/support-vector-machines-svm-an-intuitive-explanation-b084d6238106)
134 [support-vector-machines-svm-an-intuitive-explanation-b084d6238106](https://medium.com/low-code-for-advanced-data-science/support-vector-machines-svm-an-intuitive-explanation-b084d6238106).
- 135 • Liu, K. (2024). Support Vector Machines [SVM.pdf]. CPSC 8420 Advanced Machine
136 Learning, Clemson University.
- 137 • Liu, K. (2024). Principal Component Analysis [PCA.pdf]. CPSC 8420 Advanced Machine
138 Learning, Clemson University.